

FUNDAMENTALS OF MACHINE LEARNING

AA 2022-2023

Prova Finale (FACSIMILE)

9 Gennaio, 2023

Istruzioni: Niente libri, niente appunti, niente dispositivi elettronici, e niente carta per appunti. Usare matita o penna di qualsiasi colore. Usare lo spazio fornito per le risposte.

Instructions: No books, no notes, no electronic devices, and no scratch paper. Use pen or pencil. Use the space provided for your answers.

This exam has 5 questions, for a total of 100 points and 10 bonus points.

Nome: _____

Matricola: _____

1. **Multiple Choice:** Select the correct answer from the list of choices.

- (a) [5 points] True or False: A K-nearest neighbor classifier is only able to learn linear discriminant functions. ☐ True ☒ **False**
- (b) [5 points] True or False: Projecting a dataset onto its first principal component maximizes the variance of the projected data. ☒ **True** ☐ False
- (c) [5 points] True or False: The K-means algorithm is guaranteed to find the best cluster centers for any dataset. ☐ True ☒ **False**
- (d) [5 points] True or False: A Parzen kernel density estimator uses only the nearest sample in the dataset to estimate the probability of an input sample \mathbf{x} . ☐ True ☒ **False**
- (e) [5 points] How many parameters will a Multilayer Perceptron (MLP) for binary classification with a single hidden layer of width 10 and an input dimensionality of 8 have?
☐ 80 ☒ **99** ☐ 88 ☒ **None of the above** *302*
- (f) [5 points] Which of the following loss functions is called the negative log likelihood?
☐ $\mathcal{L}(\mathbf{y}, \hat{\mathbf{y}}) = -\sum_{c=1}^C (\ln y_c - \ln \hat{y}_c)^2$
☐ $\mathcal{L}(\mathbf{y}, \hat{\mathbf{y}}) = -\sum_{c=1}^C (y_c - \ln \hat{y}_c)^2$
☒ $\mathcal{L}(\mathbf{y}, \hat{\mathbf{y}}) = -\sum_{c=1}^C y_c \ln \hat{y}_c$
☐ $\mathcal{L}(\mathbf{y}, \hat{\mathbf{y}}) = -\sum_{c=1}^C \ln \hat{y}_c$
- (g) [5 points] How many iterations of gradient descent must we perform for an epoch of minibatch Stochastic Gradient Descent with a dataset of 1024 samples and a batch size of 16?
☐ 1024 ☐ 1 ☐ 32 ☒ **64**

$$N_B = \frac{1024}{16} = \frac{512}{8} = \frac{256}{4} = \frac{128}{2} = 64$$

$$\frac{512}{512} \quad \frac{256}{256} \quad \frac{128}{128} = 1 \text{ epoch}$$

se avremo anche 3 epoche

$$64 \cdot 3$$

$$(8 \cdot 10) + 10 + (10 \cdot 1) + 1$$

$$80 + 10 + 10 + 1$$

$$88 + 11 = 99$$

Total Question 1: 35

2. **Multiple Answer:** Select **ALL** correct choices: there may be more than one correct choice, but there is always at least one correct choice.

- (a) [5 points] What are the advantages of projecting data onto $K < D$ principal components?
- ☒ **We eliminate noise in the original representation.**
 - ☐ Classes are guaranteed to be linearly separable.
 - ☐ It is a nonlinear embedding that makes learning easy with simpler models.
 - ☒ **Models trained on the reduced data are simpler.**
- (b) [5 points] Which of the following are advantages of Ensemble Models (e.g. Committees)?
- ☒ **They reduce the variance of the resulting model.**
 - ☐ They are much more efficient than the base model.
 - ☒ **They can reduce the expected error of the final model.**
 - ☐ The resulting model is nonlinear even if the base model is linear.
- (c) [5 points] Which of the following are causes of the vanishing gradients when training neural networks?
- ☒ **Saturated inputs to activation functions with near-zero derivatives when saturated.**
 - ☐ Badly scaled input values.
 - ☒ **Very deep models.**
 - ☐ Bad random initialization of the network parameters.
- (d) [5 points] Which of the following are requirements for applying backpropagation to compute gradients in a deep network?
- ☐ The network must not be too deep.
 - ☒ **The network must be a directed acyclic graph.**
 - ☒ **All activation functions must be differentiable.**
 - ☐ All activation functions must be continuous.
- (e) [5 points] Which of the following are true of the Nadaraya-Watson estimator?
- ☐ It only requires some of the training data at test time.
 - ☒ **It is a nonparametric method.**
 - ☒ **It estimates a nonlinear function of the input.**
 - ☐ It estimates a linear function of the input.
- (f) [5 points] What does the learning rate control in Stochastic Gradient Descent?
- ☒ **The size of gradient steps made in each iteration.**
 - ☐ The degree of nonlinearity in the model.
 - ☐ The regression function is linear in the original input variables.
 - ☒ **The speed at which the model learns.**
- (g) [5 points] Which of the following models are nonparametric?
- ☐ The Multilayer Perceptron (MLP).
 - ☐ Logistic regression.
 - ☒ **The K-Nearest Neighbor Classifier**
 - ☐ Decision Trees.

Total Question 2: 35

3. [15 points] Show that the first principal component of dataset $\mathcal{D} = \{\mathbf{x}_n\}_{n=1}^N$ is an eigenvector of the data covariance matrix.

Solution: We saw this in class. Consider a dataset $\mathcal{D} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$ with $\mathbf{x}_n \in \mathbb{R}^D$. We seek a subspace of dimensionality $M = 1$ in which the projected points have maximum variance. The mean of the dataset \mathcal{D} projected onto a basis vector \mathbf{u}_1 is:

$$\begin{aligned} \frac{1}{N} \sum_{n=1}^N \mathbf{u}_1^T \mathbf{x}_n &= \mathbf{u}_1^T \left\{ \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n \right\} \\ &= \mathbf{u}_1^T \bar{\mathbf{x}}, \end{aligned}$$

where $\bar{\mathbf{x}} = N^{-1} \sum_n \mathbf{x}_n$. The variance is then:

$$\begin{aligned} \frac{1}{N} \sum_{n=1}^N (\mathbf{u}_1^T \mathbf{x}_n - \mathbf{u}_1^T \bar{\mathbf{x}})^2 &= \mathbf{u}_1^T \mathbf{S} \mathbf{u}_1 \\ \text{where } \mathbf{S} &= \frac{1}{N} \sum_{n=1}^N (\mathbf{x}_n - \bar{\mathbf{x}})(\mathbf{x}_n - \bar{\mathbf{x}})^T. \end{aligned}$$

We cannot simply maximize this – it is unbounded. We must constrain the optimization so that \mathbf{u} has unit norm:

$$\mathbf{u}_1^* = \arg \max_{\mathbf{u}} [\mathbf{u}_1^T \mathbf{S} \mathbf{u}_1 + \lambda_1 (1 - \mathbf{u}_1^T \mathbf{u}_1)]$$

Setting the gradient of the right-hand side to zero and solving, we obtain:

$$\begin{aligned} \nabla_{\mathbf{u}_1} \mathbf{u}_1^T \mathbf{S} \mathbf{u}_1 + \lambda_1 (1 - \mathbf{u}_1^T \mathbf{u}_1) &= \mathbf{0} \\ \implies \mathbf{S} \mathbf{u}_1 &= \lambda \mathbf{u}_1, \end{aligned}$$

which tells us that \mathbf{u}_1 must be an eigenvector of S with eigenvalue λ . □

4. [15 points] Show that a Multilayer Perceptron with two hidden layers with activation function $\sigma(x) = x$ is only capable of learning linear functions.

Solution: An MLP with two hidden layers computes the function:

$$\begin{aligned} f(\mathbf{x}) &= W_{\text{out}}\sigma(W_2\sigma(W_1\mathbf{x} + \mathbf{b}_1) + \mathbf{b}_2) + \mathbf{b}_{\text{out}} \\ &= W_{\text{out}}(W_2(W_1\mathbf{x} + \mathbf{b}_1) + \mathbf{b}_2) + \mathbf{b}_{\text{out}} \text{ (since } \sigma \text{ is the identity function)} \\ &= (W_{\text{out}}W_2W_1)\mathbf{x} + [W_{\text{out}}W_2\mathbf{b}_1 + W_{\text{out}}\mathbf{b}_2 + \mathbf{b}_{\text{out}}], \end{aligned}$$

which is a linear (well, affine) function $f(\mathbf{x}) = W\mathbf{x} + \mathbf{b}$ for:

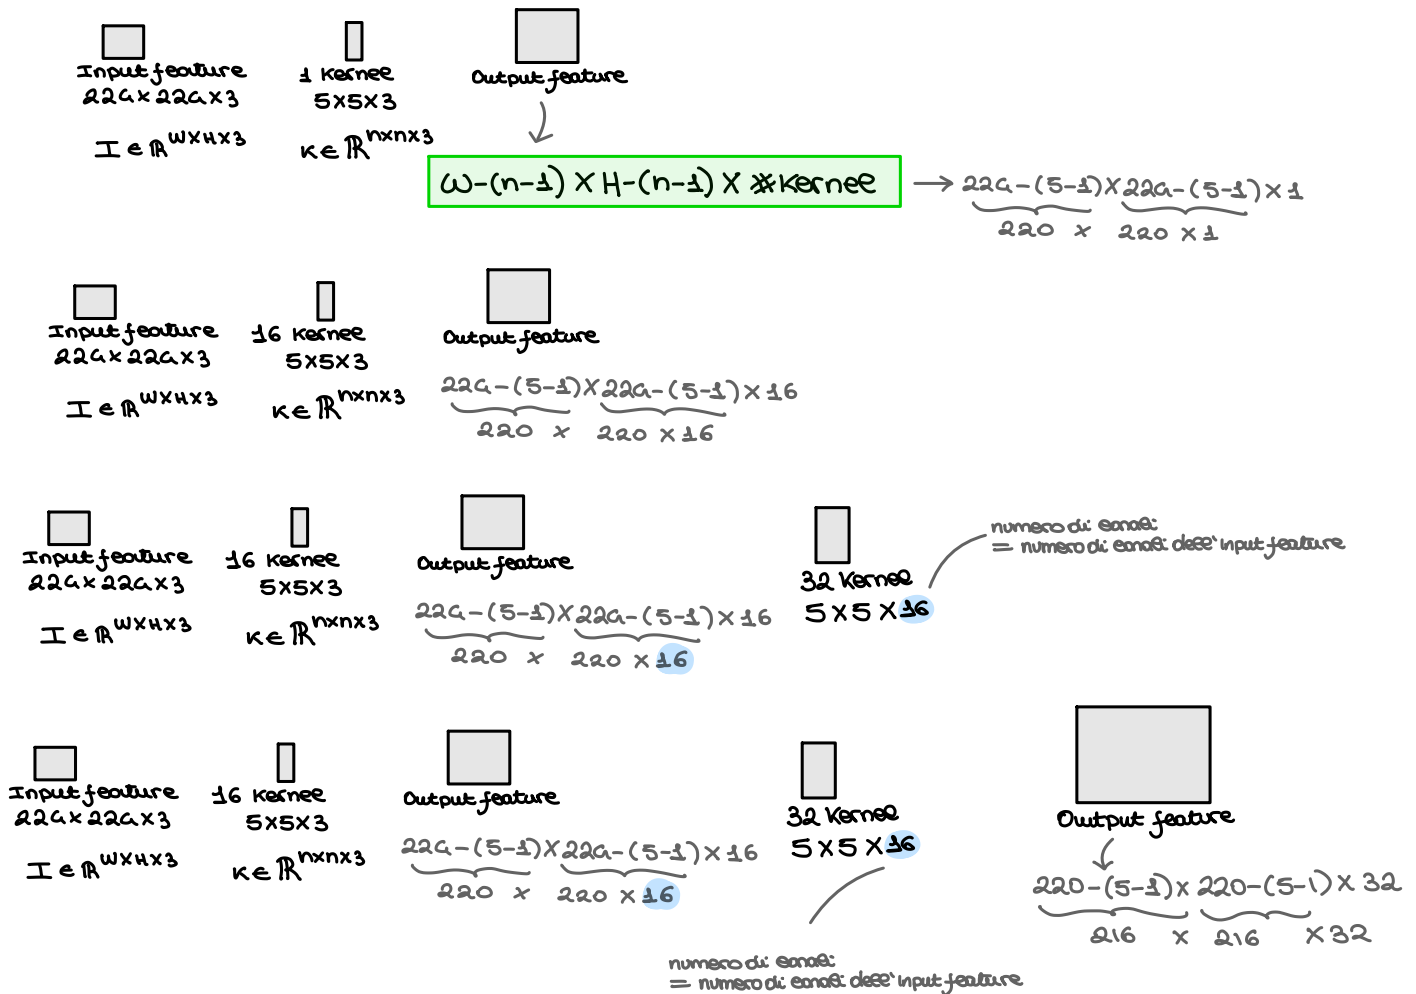
$$\begin{aligned} W &= W_{\text{out}}W_2W_1 \\ \mathbf{b} &= W_{\text{out}}W_2\mathbf{b}_1 + W_{\text{out}}\mathbf{b}_2 + \mathbf{b}_{\text{out}}. \end{aligned}$$

□

5. [10 points (bonus)] Design a Deep Convolutional Neural Network (with at least three convolutional layers and one or more pooling layers) to classify MNIST images (input size 28×28). Draw the network (or write pseudocode for its definition) and indicate how many parameters each layer has and the sizes of the intermediate feature maps.

Solution: I will write pseudocode in tabular form for the definition of each layer (with corresponding numbers of parameters and size of the activations:

Layer	Type	Activation Size	# Parameters
1	Input	$1 \times 28 \times 28$	0
2	Conv2D(32, 1, 3, 3)	$32 \times 26 \times 26$	320 ($32 * 3 * 3 + 32$)
3	ReLU	$32 \times 26 \times 26$	0
4	Conv2D(32, 32, 3, 3)	$32 \times 24 \times 24$	9248
5	ReLU	$32 \times 24 \times 24$	0
6	MaxPool(2, 2)	$32 \times 13 \times 13$	0
7	Conv2D(16, 32, 3, 3)	$16 \times 11 \times 11$	4624
8	ReLU	$16 \times 11 \times 11$	0
9	Conv2D(16, 16, 3, 3)	$16 \times 9 \times 9$	2320
10	ReLU	$16 \times 9 \times 9$	0
11	MaxPool(2, 2)	$16 \times 5 \times 5$	0
12	Flatten()	400	0
13	Linear(400, 128)	128	51328
14	ReLU	128	0
15	Linear(128, 64)	64	8256
16	ReLU	64	0
17	Linear(64, 10)	10	650



5. [10 points (bonus)] Design a Deep Convolutional Neural Network (with at least three convolutional layers and one or more pooling layers) to classify MNIST images (input size 28×28). Draw the network (or write pseudocode for its definition) and indicate how many parameters each layer has and the sizes of the intermediate feature maps.

Solution: I will write pseudocode in tabular form for the definition of each layer (with corresponding numbers of parameters and size of the activations):

Layer	Type	Activation Size	# Parameters
1	Input	$1 \times 28 \times 28$	0
2	Conv2D(32, 1, 3, 3)	$32 \times 26 \times 26$	320 ($32 * 3 * 3 + 32$)
3	ReLU	$32 \times 26 \times 26$	0
4	Conv2D(32, 32, 3, 3)	$32 \times 24 \times 24$	9248
5	ReLU	$32 \times 26 \times 26$	0
6	MaxPool(2, 2)	$32 \times 13 \times 13$	0
7	Conv2D(16, 32, 3, 3)	$16 \times 11 \times 11$	4624
8	ReLU	$16 \times 11 \times 11$	0
9	Conv2D(16, 16, 3, 3)	$16 \times 9 \times 9$	2320
10	ReLU	$16 \times 9 \times 9$	0
11	MaxPool(2, 2)	$16 \times 5 \times 5$	0
12	Flatten()	400	0
13	Linear(400, 128)	128	51328
14	ReLU	128	0
15	Linear(128, 64)	64	8256
16	ReLU	64	0
17	Linear(64, 10)	10	650

1. INPUT	$28 \times 28 \times 1$	0
2. CONV2D(32, 1, 3, 3)	$26 \times 26 \times 32$	$32 \times 1 \times 3 \times 3 + 32$
3. ReLU	$26 \times 26 \times 32$	0
4. CONV2D(32, 32, 3, 3)	$26 - (3 - 1) \times 26 - (3 - 1) \times 32$	$32 \times 32 \times 3 \times 3 + 32$
5. ReLU	$24 \times 24 \times 32$	0
6. MaxPool(2, 2)	$12 \times 12 \times 32$	0
7. CONV2D(16, 32, 3, 3)	$12 - (3 - 1) \times 10 \times 10 \times 16$	$16 \times 32 \times 3 \times 3 + 16$
8. ReLU	$10 \times 10 \times 16$	0
9. FLATTEN()	400	0
10. LINEAR(400, 128)	128	$400 \times 128 + 128$
11. ReLU	128	0
12. LINEAR(128, 64)	64	$128 \times 64 + 64$
13. ReLU	64	0
14. LINEAR(64, 10)	10	$64 \times 10 + 10$

